

H3: Accuracy vs. latency — Pareto frontier & baseline reference (fp32 = GPU inference, int8 = CPU inference; per-precision deployment target)

RTX 4060 Laptop (local)

