

RTX 4060 Laptop (local): layer latency per FLOP vs. kernel size (dense/groups=1, res=64)

